

```
install.packages("magrittr") install.packages("ggmap") install.packages("geosphere") install.packages("httr")
```

```
In [1]: install.packages("rpart.plot")  
install.packages("inTrees")
```

```
Updating HTML index of packages in '.Library'  
Making 'packages.html' ... done  
also installing the dependencies 'RRF', 'arules', 'gbm'
```

```
Warning message in install.packages("inTrees"):  
"installation of package 'RRF' had non-zero exit status"Warning message  
in install.packages("inTrees"):  
"installation of package 'arules' had non-zero exit status"Warning mess  
age in install.packages("inTrees"):  
"installation of package 'gbm' had non-zero exit status"Warning message  
in install.packages("inTrees"):  
"installation of package 'inTrees' had non-zero exit status"Updating HT  
ML index of packages in '.Library'  
Making 'packages.html' ... done
```

King County Housing Prices

Sruthi Jogi, Jonathan McFadden, Angela Zhao

TCSS-551: Big Data Analytics :- Autumn 2017

Introduction

Overview

For our final project, we have chosen to analyze data covering housing sales in King County. To do this, we are using the data from the **Kaggle** *King County House Sales Prediction* page at <https://www.kaggle.com/harlfoxem/housesalesprediction>

From this page, we sign-up for an account (*free, but required for downloading*) and then download the *zip* file containing the CSV file with the data.

Our goal is to use this data to create models for home sales in King County based on the feature information provided in the obtained data file. Our eventual goal is two-fold. First, we wish to create a model or models which will enable us to quantitatively predict house sale prices, using this data set as the basis for our model or models. Our other goal is to determine, based on the obtained data, which features are most important to the sale price of a house.

Data File

Our first task is to import, examine, and then give an overall description of the data. We are especially interested in the size and descriptive contents of the data file. Specifically, we want to know the number of sales contained within the data file and, especially, what parameters the data file uses to describe each house sale. Furthermore, we want to check the import to ensure that the data was initially complete, that it was then imported correctly, and that **R** is interpreting the imported data properly.

Import and First-Look

We begin by importing the data file into the '**houseDFo**()' data frame. This data frame will serve as an initial data-frame, not the working one. This is because we may need an initial frame to reload as we clean the data, allowing us to avoid having to reimport the CSV file over and over again. Thus, we now import the CSV file into this initial data frame.

```
In [2]: houseDFo <- read.csv("../houseData.csv")
```

We are now interested in the number of data-points contained within the data file. Thus, we want to see how many rows **R** has imported.

```
In [3]: nrow(houseDFo)
```

```
21613
```

We also want to see how many descriptors the imported data uses to describe each house sale. Thus we want to see how many columns **R** has imported.

```
In [4]: ncol(houseDFo)
```

```
21
```

In addition, we want to see what the labels for those columns are and what type of values the elements of each column have (*integer, numeric, string, etc.*)

```
In [5]: sapply(houseDFo, class)
```

id	'numeric'
date	'factor'
price	'numeric'
bedrooms	'integer'
bathrooms	'numeric'
sqft_living	'integer'
sqft_lot	'integer'
floors	'numeric'
waterfront	'integer'
view	'integer'
condition	'integer'
grade	'integer'
sqft_above	'integer'
sqft_basement	'integer'
yr_built	'integer'
yr_renovated	'integer'
zipcode	'integer'
lat	'numeric'
long	'numeric'
sqft_living15	'integer'
sqft_lot15	'integer'

From above, it is clear that the **date** column did not import as a *date*, instead importing as a *factor*. Therefore, we will now examine the first few rows of the imported data to see what may have caused the issues with importation.

```
In [6]: head(houseDfO)
```

id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floor
7129300520	20141013T000000	221900	3	1.00	1180	5650	1
6414100192	20141209T000000	538000	3	2.25	2570	7242	2
5631500400	20150225T000000	180000	2	1.00	770	10000	1
2487200875	20141209T000000	604000	4	3.00	1960	5000	1
1954400510	20150218T000000	510000	3	2.00	1680	8080	1
7237550310	20140512T000000	1225000	4	4.50	5420	101930	1

Clearly, some elements of the data file did not import correctly; therefore, we must clean the data before we can proceed to analysis.

Clean the Data

Missing Data

First, we will check to see if there are any missing data points.

```
In [7]: houseDfO[!complete.cases(houseDfO),]
```

```
Warning message in cbind(parts$left, ellip_h, parts$right, deparse.level = 0L):
```

```
"number of rows of result is not a multiple of vector length (arg 2)"Warning message in cbind(parts$left, ellip_h, parts$right, deparse.level = 0L):
```

```
"number of rows of result is not a multiple of vector length (arg 2)"Warning message in cbind(parts$left, ellip_h, parts$right, deparse.level = 0L):
```

```
"number of rows of result is not a multiple of vector length (arg 2)"Warning message in cbind(parts$left, ellip_h, parts$right, deparse.level = 0L):
```

```
"number of rows of result is not a multiple of vector length (arg 2)"
```

id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	gr
----	------	-------	----------	-----------	-------------	----------	--------	------------	------	-----	----

Since there are no missing data points, we can move on to the dates.

Dates

From the first few rows of the data table seen above, it is clear that we must first strip the "T000000" string at the end of every date. To do this, we require the **stringr** library. Thus, we import **stringr**

```
In [8]: library(stringr)
```

so we can now strip the offending substrings. Before stripping these substrings, we create a copy of our initial data frame, **houseDFo()**, so that our initial import data frame will remain untouched, and therefore available for reloading other data frames. Thus, we create the copy and strip the substrings, storing the result in the copied data frame **houseDFo1()**.

```
In [9]: houseDFo1 <- houseDFo
        houseDFo1$date = str_replace(houseDFo$date, "T000000", "")
```

We now examine the result of this

```
In [10]: head(houseDFo1)
```

id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
7129300520	20141013	221900	3	1.00	1180	5650	1	0
6414100192	20141209	538000	3	2.25	2570	7242	2	0
5631500400	20150225	180000	2	1.00	770	10000	1	0
2487200875	20141209	604000	4	3.00	1960	5000	1	0
1954400510	20150218	510000	3	2.00	1680	8080	1	0
7237550310	20140512	1225000	4	4.50	5420	101930	1	0

The dates are now just strings of numbers with the format 'yyyymmdd'; therefore, we can use the date conversion method from **R** to convert these dates.

```
In [11]: houseDFo1 <- transform(houseDFo1, date = as.Date(date, "%Y%m%d"))
```

To ensure that the conversion to dates happend properly, we will no check the column data types followed by looking at the first few rows of the data.

```
In [12]: sapply(houseDFo1, class)
         head(houseDFo1)
```

```

      id 'numeric'
     date 'Date'
    price 'numeric'
  bedrooms 'integer'
  bathrooms 'numeric'
 sqft_living 'integer'
 sqft_lot 'integer'
    floors 'numeric'
  waterfront 'integer'
    view 'integer'
   condition 'integer'
    grade 'integer'
 sqft_above 'integer'
sqft_basement 'integer'
    yr_built 'integer'
 yr_renovated 'integer'
    zipcode 'integer'
    lat 'numeric'
    long 'numeric'
sqft_living15 'integer'
 sqft_lot15 'integer'

```

id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
7129300520	2014-10-13	221900	3	1.00	1180	5650	1	0
6414100192	2014-12-09	538000	3	2.25	2570	7242	2	0
5631500400	2015-02-25	180000	2	1.00	770	10000	1	0
2487200875	2014-12-09	604000	4	3.00	1960	5000	1	0
1954400510	2015-02-18	510000	3	2.00	1680	8080	1	0
7237550310	2014-05-12	1225000	4	4.50	5420	101930	1	0

Since the results for the date conversions are as desired, we can now store the data in a final data frame followed by moving on to beginning our analysis.

```
In [13]: houseDF <- houseDFo1
```

We will also create a version of the data with the **ID** column stripped out.

```
In [14]: houseDFa <- houseDF[-c(1)]
```

Initial Analysis

To begin our analysis, we will look at the basic statistics of every column (*except the date*).

```
In [15]: sapply(houseDFa[-c(1)], function(x) list(mean=mean(x),  
                                                stdev=sd(x, na.rm=TRUE)))
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
mean	540088.1	3.370842	2.114757	2079.9	15106.97	1.494309	0.007541757	0.2
stdev	367127.2	0.9300618	0.7701632	918.4409	41420.51	0.5399889	0.0865172	0.7

and get a summary of the entire

```
In [16]: summary(houseDFa)
```

date		price		bedrooms		bathrooms	
Min.	:2014-05-02	Min.	: 75000	Min.	: 0.000	Min.	:0.000
1st Qu.:	2014-07-22	1st Qu.:	321950	1st Qu.:	3.000	1st Qu.:	1.750
Median	:2014-10-16	Median	: 450000	Median	: 3.000	Median	:2.250
Mean	:2014-10-29	Mean	: 540088	Mean	: 3.371	Mean	:2.115
3rd Qu.:	2015-02-17	3rd Qu.:	645000	3rd Qu.:	4.000	3rd Qu.:	2.500
Max.	:2015-05-27	Max.	:7700000	Max.	:33.000	Max.	:8.000

sqft_living		sqft_lot		floors		waterfront	
Min.	: 290	Min.	: 520	Min.	:1.000	Min.	:0.000000
1st Qu.:	1427	1st Qu.:	5040	1st Qu.:	1.000	1st Qu.:	0.000000
Median	: 1910	Median	: 7618	Median	:1.500	Median	:0.000000
Mean	: 2080	Mean	: 15107	Mean	:1.494	Mean	:0.007542
3rd Qu.:	2550	3rd Qu.:	10688	3rd Qu.:	2.000	3rd Qu.:	0.000000
Max.	:13540	Max.	:1651359	Max.	:3.500	Max.	:1.000000

view		condition		grade		sqft_above	
Min.	:0.0000	Min.	:1.000	Min.	: 1.000	Min.	: 290
1st Qu.:	0.0000	1st Qu.:	3.000	1st Qu.:	7.000	1st Qu.:	1190
Median	:0.0000	Median	:3.000	Median	: 7.000	Median	:1560
Mean	:0.2343	Mean	:3.409	Mean	: 7.657	Mean	:1788
3rd Qu.:	0.0000	3rd Qu.:	4.000	3rd Qu.:	8.000	3rd Qu.:	2210
Max.	:4.0000	Max.	:5.000	Max.	:13.000	Max.	:9410

sqft_basement		yr_built		yr_renovated		zipcode	
Min.	: 0.0	Min.	:1900	Min.	: 0.0	Min.	:98001
1st Qu.:	0.0	1st Qu.:	1951	1st Qu.:	0.0	1st Qu.:	98033
Median	: 0.0	Median	:1975	Median	: 0.0	Median	:98065
Mean	: 291.5	Mean	:1971	Mean	: 84.4	Mean	:98078
3rd Qu.:	560.0	3rd Qu.:	1997	3rd Qu.:	0.0	3rd Qu.:	98118
Max.	:4820.0	Max.	:2015	Max.	:2015.0	Max.	:98199

lat		long		sqft_living15		sqft_lot15	
Min.	:47.16	Min.	: -122.5	Min.	: 399	Min.	: 651
1st Qu.:	47.47	1st Qu.:	-122.3	1st Qu.:	1490	1st Qu.:	5100
Median	:47.57	Median	: -122.2	Median	:1840	Median	: 7620
Mean	:47.56	Mean	: -122.2	Mean	:1987	Mean	: 12768
3rd Qu.:	47.68	3rd Qu.:	-122.1	3rd Qu.:	2360	3rd Qu.:	10083
Max.	:47.78	Max.	: -121.3	Max.	:6210	Max.	:871200

Model A

We also run a simple linear model on the *entire* dataset so that we can see how significant each variable is to determining the price (*basically running a t-Test on all variables*). To do this, we need to **nnet** library, so we load it

```
In [17]: library(nnet)
```


Then we run the model, our first, and display the results.

```
In [18]: modelA <- lm(price ~., data=houseDFa)
summary(modelA)
```

```
Call:
lm(formula = price ~ ., data = houseDFa)

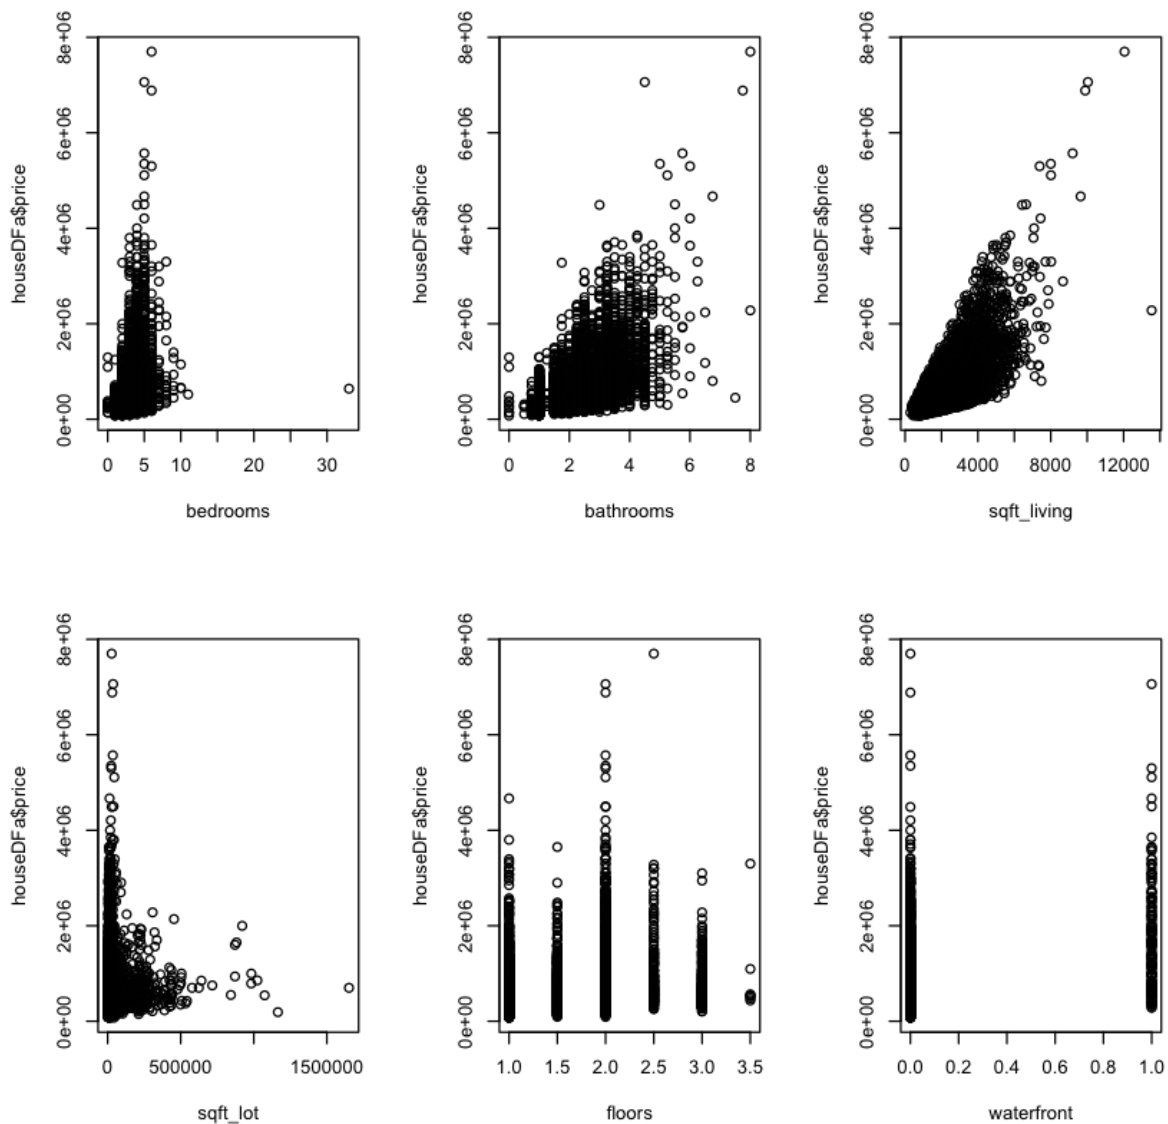
Residuals:
    Min       1Q   Median       3Q      Max
-1306672  -98900   -8963    77327  4330103

Coefficients: (1 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.618e+06  2.933e+06   1.574  0.11539
date         1.165e+02  1.213e+01   9.608 < 2e-16 ***
bedrooms     -3.588e+04  1.888e+03 -19.005 < 2e-16 ***
bathrooms     4.137e+04  3.247e+03  12.741 < 2e-16 ***
sqft_living   1.502e+02  4.376e+00  34.327 < 2e-16 ***
sqft_lot      1.257e-01  4.782e-02   2.629  0.00858 **
floors        7.158e+03  3.589e+03   1.995  0.04610 *
waterfront    5.826e+05  1.732e+04  33.628 < 2e-16 ***
view          5.260e+04  2.136e+03  24.629 < 2e-16 ***
condition     2.774e+04  2.351e+03  11.799 < 2e-16 ***
grade         9.624e+04  2.149e+03  44.791 < 2e-16 ***
sqft_above    3.084e+01  4.351e+00   7.088 1.40e-12 ***
sqft_basement      NA           NA      NA      NA
yr_built      -2.618e+03  7.251e+01 -36.113 < 2e-16 ***
yr_renovated   2.079e+01  3.649e+00   5.698 1.23e-08 ***
zipcode       -5.807e+02  3.292e+01 -17.643 < 2e-16 ***
lat           6.053e+05  1.072e+04  56.487 < 2e-16 ***
long          -2.136e+05  1.311e+04 -16.300 < 2e-16 ***
sqft_living15  2.195e+01  3.441e+00   6.381 1.79e-10 ***
sqft_lot15     -3.825e-01  7.311e-02  -5.232 1.69e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

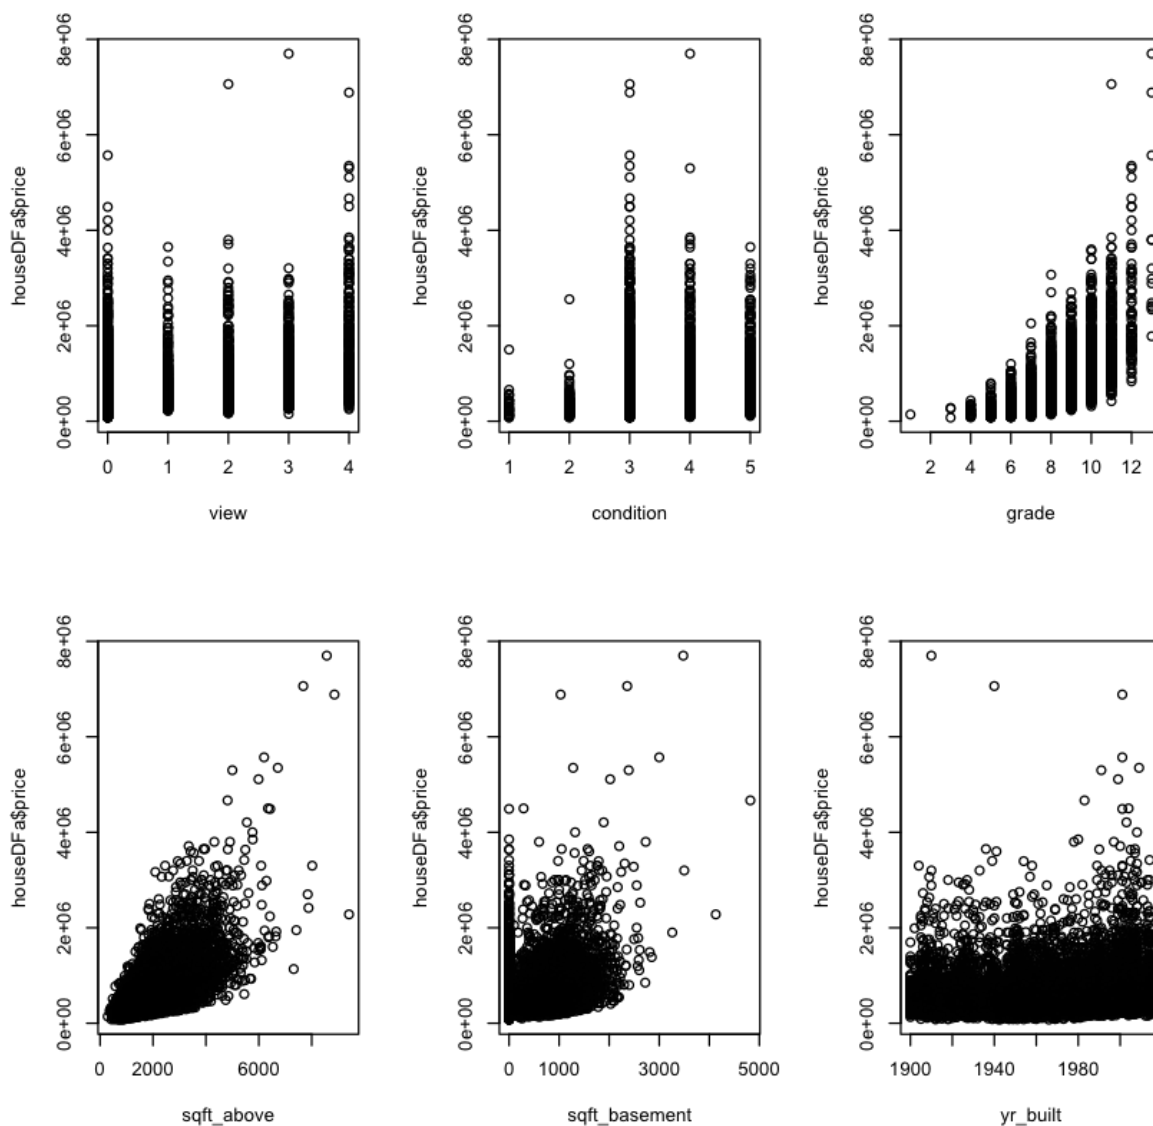
Residual standard error: 200800 on 21594 degrees of freedom
Multiple R-squared:  0.701,    Adjusted R-squared:  0.7008
F-statistic: 2813 on 18 and 21594 DF,  p-value: < 2.2e-16
```

We also plot the data to look for potential relationships and outliers.

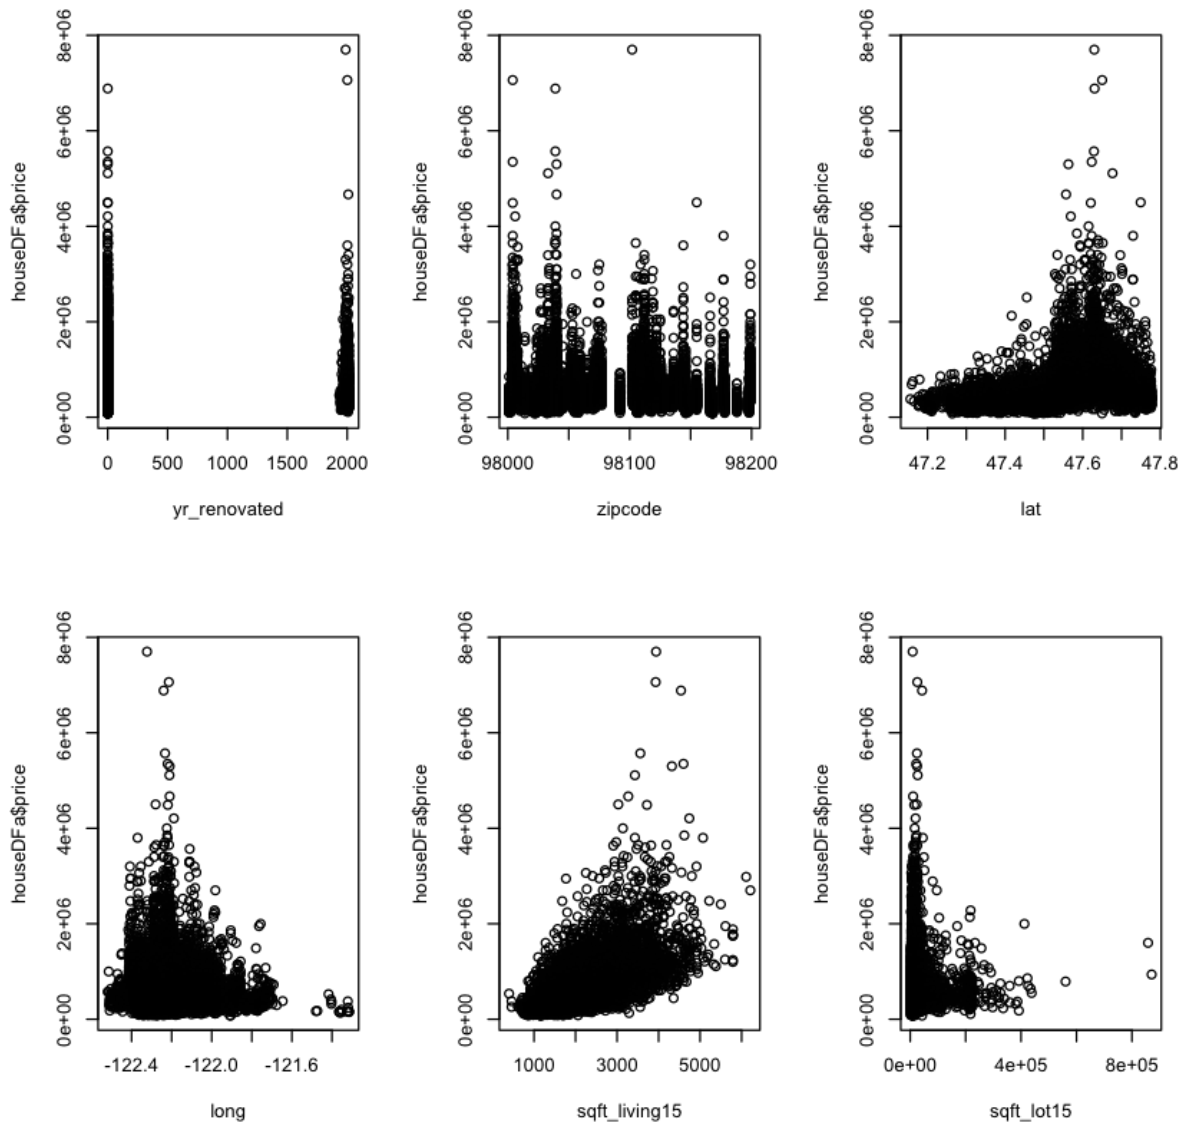
```
In [19]: par(mfrow=c(2,3))  
for(i in 3:8) {plot(houseDf[,i], houseDf$price, xlab=names(houseDf[,i]),  
ylab=names(houseDf$price))}
```



```
In [20]: par(mfrow=c(2,3))
for(i in 9:14) {plot(houseDf[,i], houseDf$price, xlab=names(houseDf[i
]), ylab=names(houseDf$price))}
```



```
In [21]: par(mfrow=c(2,3))
for(i in 15:20) {plot(houseDfa[,i], houseDfa$price, xlab=names(houseDfa[
i]), ylab=names(houseDfa$price))}
```



Model B

For our second model, we chose a decision tree, to determine which home characteristics were most important, followed by a random forest to provide a more accurate regression than what a lone decision tree would provide. First, we need the decision tree library

```
In [22]: library(rpart)
```

and then we create the decision tree

```

In [23]: modelB1 <- rpart(houseDf$price ~., houseDf[, -c(1,2)], na.action = na.r
part)
modelB1

n= 21613

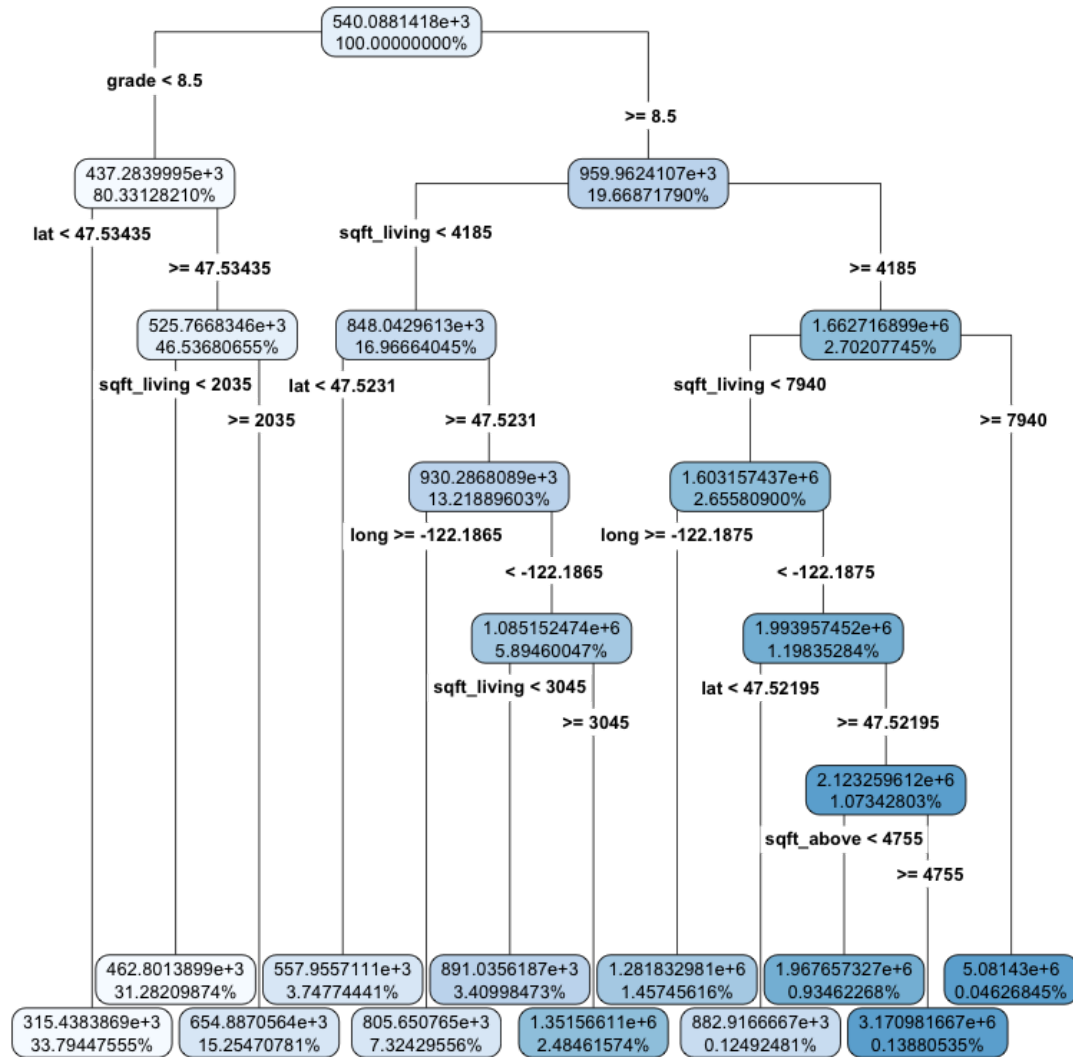
node), split, n, deviance, yval
* denotes terminal node

1) root 21613 2.912917e+15 540088.1
 2) grade< 8.5 17362 6.656597e+14 437284.0
   4) lat< 47.53435 7304 1.003407e+14 315438.4 *
   5) lat>=47.53435 10058 3.781350e+14 525766.8
      10) sqft_living< 2035 6761 1.361306e+14 462801.4 *
      11) sqft_living>=2035 3297 1.602317e+14 654887.1 *
 3) grade>=8.5 4251 1.314336e+15 959962.4
   6) sqft_living< 4185 3667 4.934437e+14 848043.0
      12) lat< 47.5231 810 3.289091e+13 557955.7 *
      13) lat>=47.5231 2857 3.730659e+14 930286.8
          26) long>=-122.1865 1583 6.476840e+13 805650.8 *
          27) long< -122.1865 1274 2.531521e+14 1085152.0
              54) sqft_living< 3045 737 7.221140e+13 891035.6 *
              55) sqft_living>=3045 537 1.150553e+14 1351566.0 *
 7) sqft_living>=4185 584 4.865430e+14 1662717.0
   14) sqft_living< 7940 574 3.372812e+14 1603157.0
      28) long>=-122.1875 315 7.499599e+13 1281833.0 *
      29) long< -122.1875 259 1.902060e+14 1993957.0
          58) lat< 47.52195 27 6.625440e+12 882916.7 *
          59) lat>=47.52195 232 1.463726e+14 2123260.0
              118) sqft_above< 4755 202 8.983921e+13 1967657.0 *
              119) sqft_above>=4755 30 1.871090e+13 3170982.0 *
 15) sqft_living>=7940 10 3.034967e+13 5081430.0 *


```

and plot it

```
In [24]: library(rpart.plot)
rpart.plot(modelB1,digits=10,fallen.leaves=TRUE,type=4)
```



```
In [25]: plotcp(modelB1)
```



We now load the random forest library

```
In [26]: library(randomForest)
```

randomForest 4.6-12

Type rfNews() to see new features/changes/bug fixes.

and then create the random forest model for our data

```
In [27]: modelB2 <- randomForest(houseDfa$price ~., houseDfa[,-c(1,2)])
modelB2


Call:
  randomForest(formula = houseDfa$price ~ ., data = houseDfa[,      -c
(1, 2)])

      Type of random forest: regression
    Number of trees: 500
No. of variables tried at each split: 6

      Mean of squared residuals: 16114715551
        % Var explained: 88.04
```

which we then plot

```
In [28]: plot(modelB2)
```



```
In [29]: #partialPlot(modelB2, houseDfa[,-c(1,2)], houseDfa$bathrooms)
```


We also want to use our random forest to see which variables are more important when using an ensemble approach. To do this, we create a table of the importance of each variable in the random forest

```
In [31]: importance(modelB2)
```

	IncNodePurity
bedrooms	1.557039e+13
bathrooms	1.176603e+14
sqft_living	6.459734e+14
sqft_lot	4.835966e+13
floors	1.028707e+13
waterfront	7.023169e+13
view	7.662750e+13
condition	1.323093e+13
grade	5.707538e+14
sqft_above	2.201358e+14
sqft_basement	5.312903e+13
yr_built	9.719926e+13
yr_renovated	9.267914e+12
zipcode	7.040288e+13
lat	4.021602e+14
long	1.697445e+14
sqft_living15	2.490493e+14
sqft_lot15	5.004354e+13

Model C

Our final model was a logistic regression. We chose this model because we were hoping to find any non-linear relationship that might be in play between various properties of a house and its price. In constructing this model, we assumed that, overall, the parameters of the houses were gaussian.

```
In [38]: modelC <- glm(houseDFa$price ~., houseDFa[,-c(1,2)], family = gaussian(link = "identity"))
modelC
summary(modelC)
```

```
Call: glm(formula = houseDFa$price ~ ., family = gaussian(link = "identity"),
  data = houseDFa[, -c(1, 2)])
```

Coefficients:

(Intercept)	bedrooms	bathrooms	sqft_living	sqft_lot
6.690e+06	-3.577e+04	4.114e+04	1.501e+02	1.286e-01
floors	waterfront	view	condition	grade
6.690e+03	5.830e+05	5.287e+04	2.639e+04	9.589e+04
sqft_above	sqft_basement	yr_built	yr_renovated	zipcode
3.113e+01	NA	-2.620e+03	1.981e+01	-5.824e+02
lat	long	sqft_living15	sqft_lot15	
6.027e+05	-2.147e+05	2.168e+01	-3.826e-01	

Degrees of Freedom: 21612 Total (i.e. Null); 21595 Residual

Null Deviance: 2.913e+15

Residual Deviance: 8.746e+14 AIC: 589200

```
Call:
glm(formula = houseDFa$price ~ ., family = gaussian(link = "identity"),
    data = houseDFa[, -c(1, 2)])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1291725	-99229	-9739	77583	4333222

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	6.690e+06	2.931e+06	2.282	0.02249	*
bedrooms	-3.577e+04	1.892e+03	-18.906	< 2e-16	***
bathrooms	4.114e+04	3.254e+03	12.645	< 2e-16	***
sqft_living	1.501e+02	4.385e+00	34.227	< 2e-16	***
sqft_lot	1.286e-01	4.792e-02	2.683	0.00729	**
floors	6.690e+03	3.596e+03	1.860	0.06285	.
waterfront	5.830e+05	1.736e+04	33.580	< 2e-16	***
view	5.287e+04	2.140e+03	24.705	< 2e-16	***
condition	2.639e+04	2.351e+03	11.221	< 2e-16	***
grade	9.589e+04	2.153e+03	44.542	< 2e-16	***
sqft_above	3.113e+01	4.360e+00	7.139	9.71e-13	***
sqft_basement	NA	NA	NA	NA	
yr_built	-2.620e+03	7.266e+01	-36.062	< 2e-16	***
yr_renovated	1.981e+01	3.656e+00	5.420	6.03e-08	***
zipcode	-5.824e+02	3.299e+01	-17.657	< 2e-16	***
lat	6.027e+05	1.073e+04	56.149	< 2e-16	***
long	-2.147e+05	1.313e+04	-16.349	< 2e-16	***
sqft_living15	2.168e+01	3.448e+00	6.289	3.26e-10	***
sqft_lot15	-3.826e-01	7.327e-02	-5.222	1.78e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 40500645795)

Null deviance: 2.9129e+15 on 21612 degrees of freedom
 Residual deviance: 8.7461e+14 on 21595 degrees of freedom
 AIC: 589244

Number of Fisher Scoring iterations: 2

Summary

We discovered that the most important properties of a house, with respect to influence on its price, were

- The number of bedrooms
- The number of bathrooms
- The condition of the house
- Whether the house had a view or was on waterfront property
- The amount of living space in the house, as measured in sq ft

We also found that the random forest was the most accurate method, with each method having accuracies of

- Linear regression: ~71%
- Random forest: ~88%
- Non-Linear Logistic Regression: *no accuracy* but residuals were higher than the **Random Forest** model.

Finally, because there are clearly some non-linear dynamics at play in the data, we believe it would be useful to have information on what **type** each house is. Is it a condominium? A "*regular*" house? A townhome? We would also like data on the quality of things near the house like schools, parks, transit, and roads. These may be factors which could be either playing an unseen role in determining the price of a house **or** influencing other variables in such a manner as to make them appear non-linear. With this information, we might be able to linearize those variables.